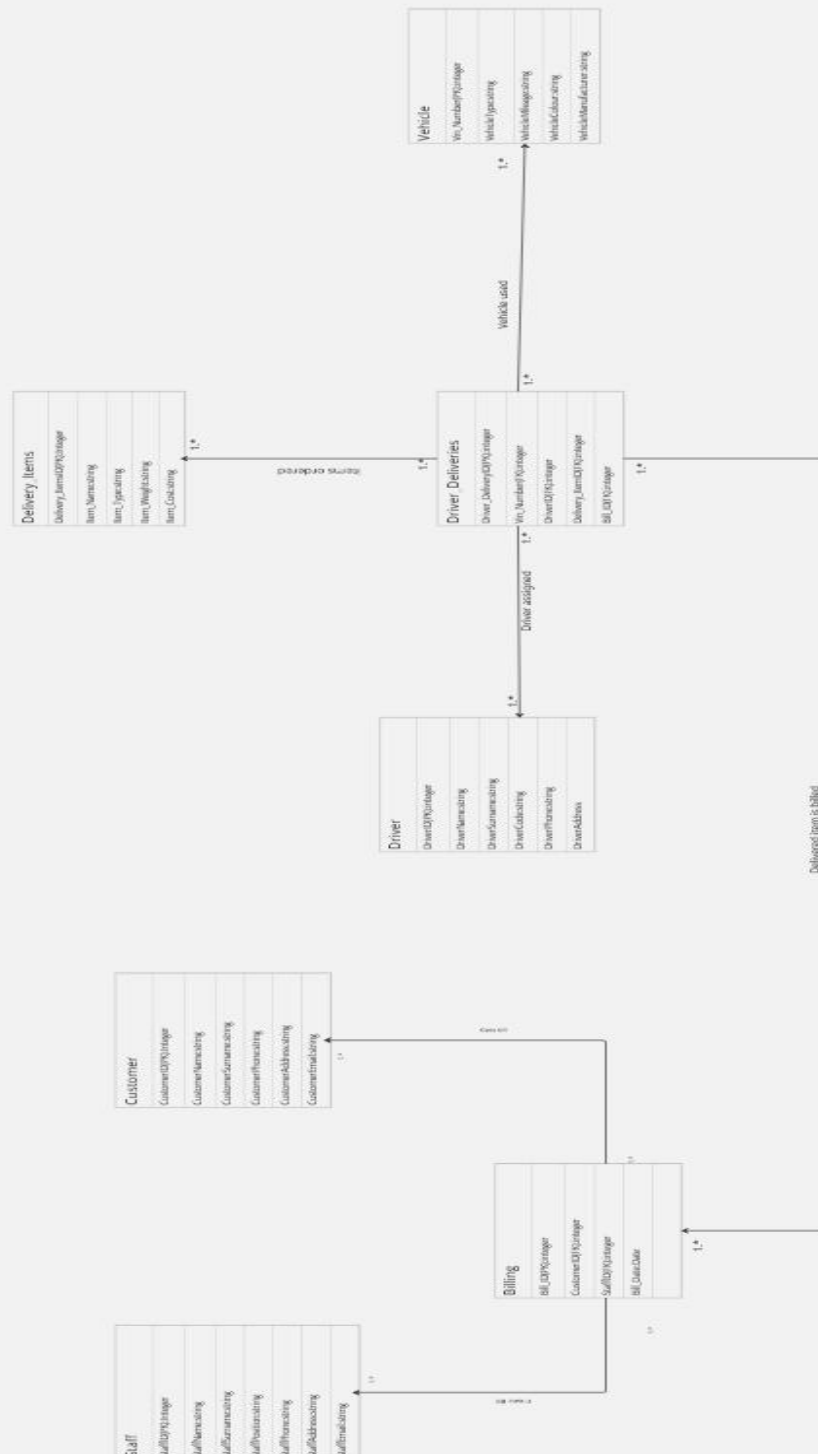




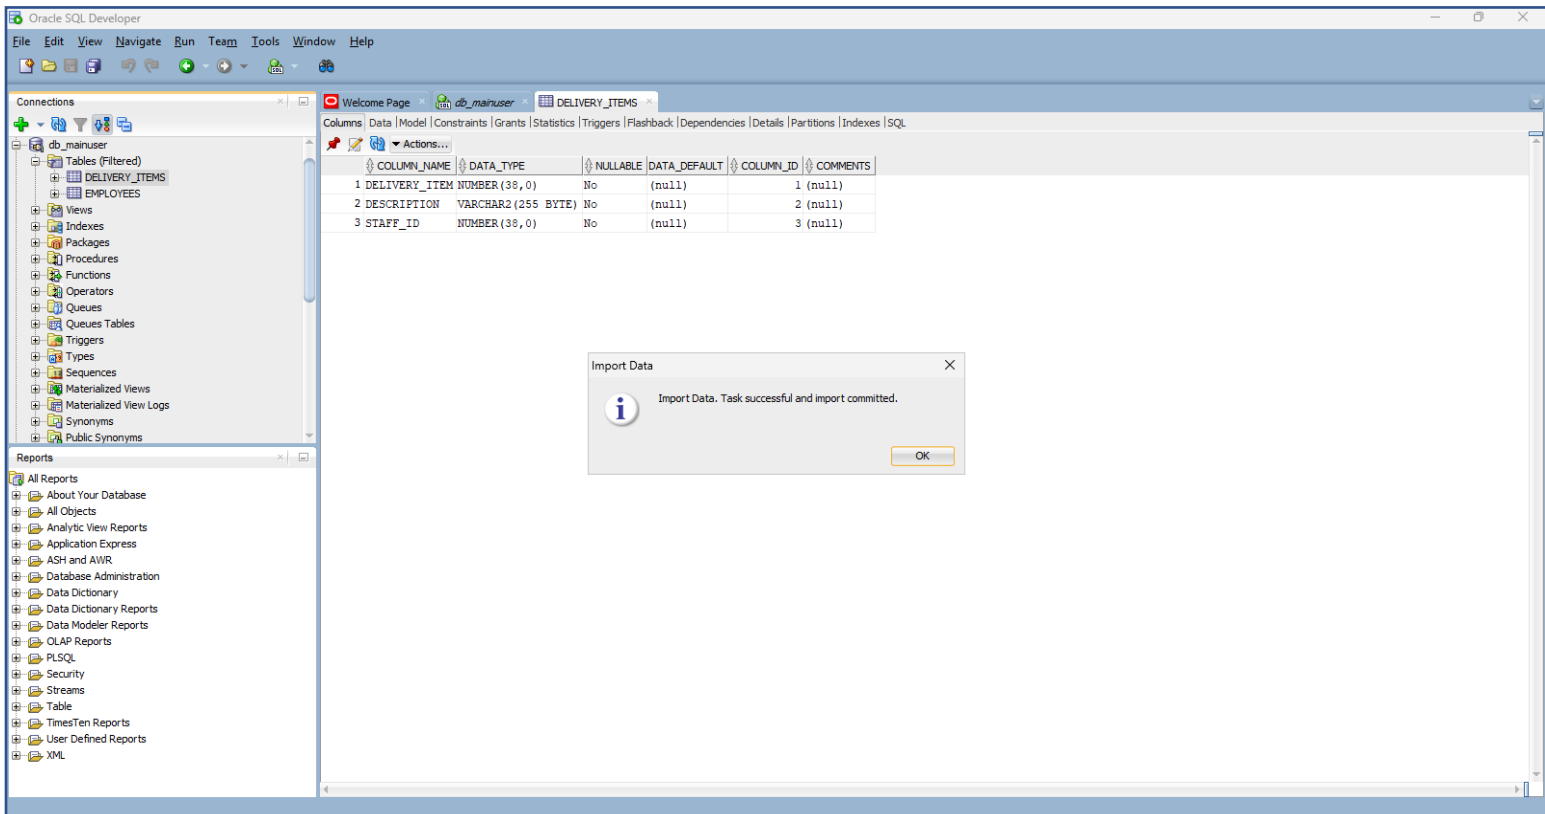
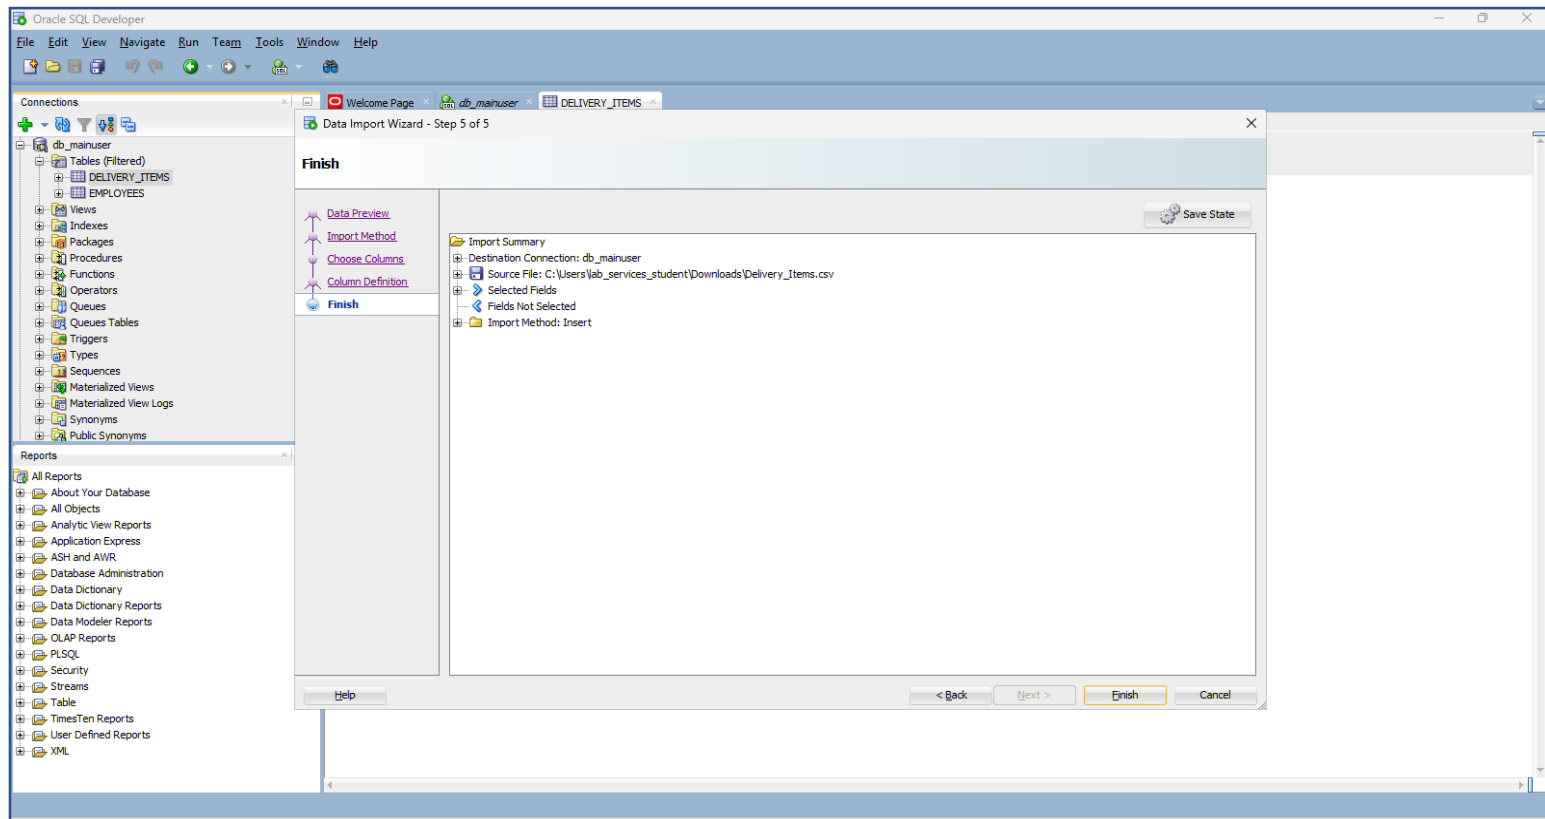
ST10255824-ASSIGNMENT 1

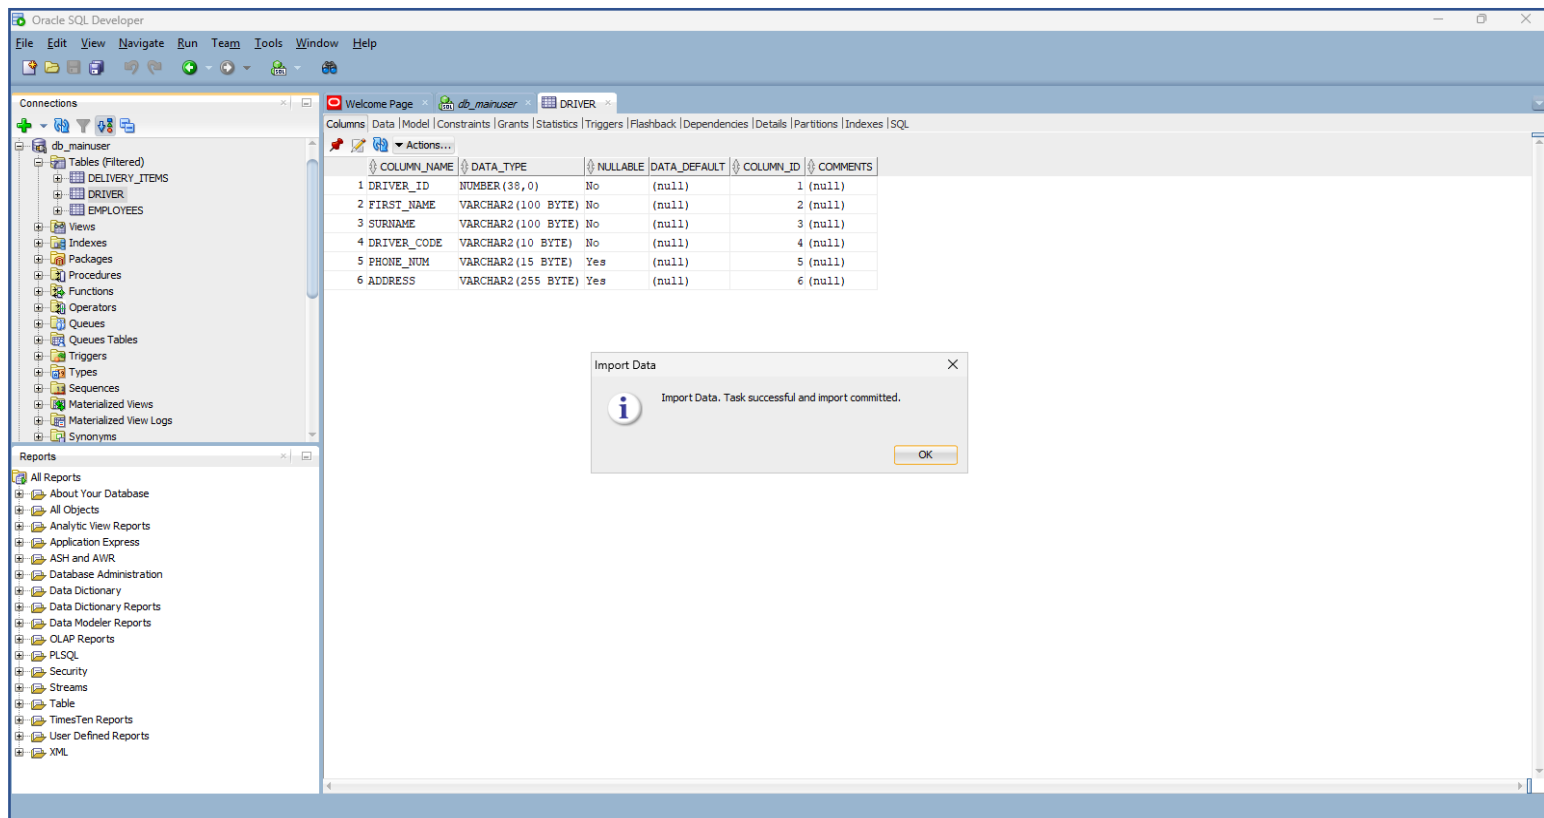

LCHATEAU-INSY7213

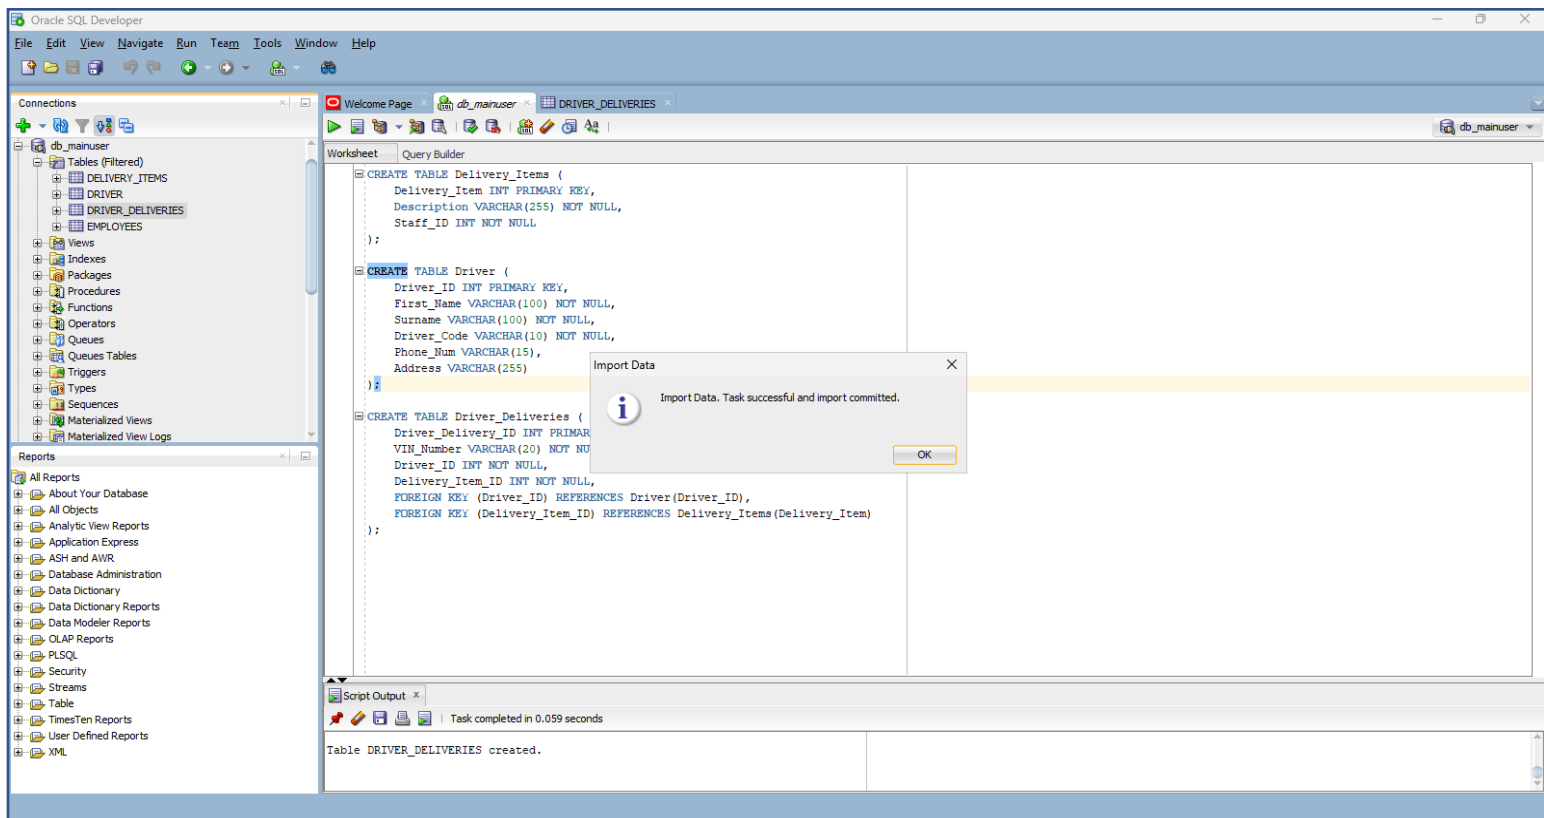
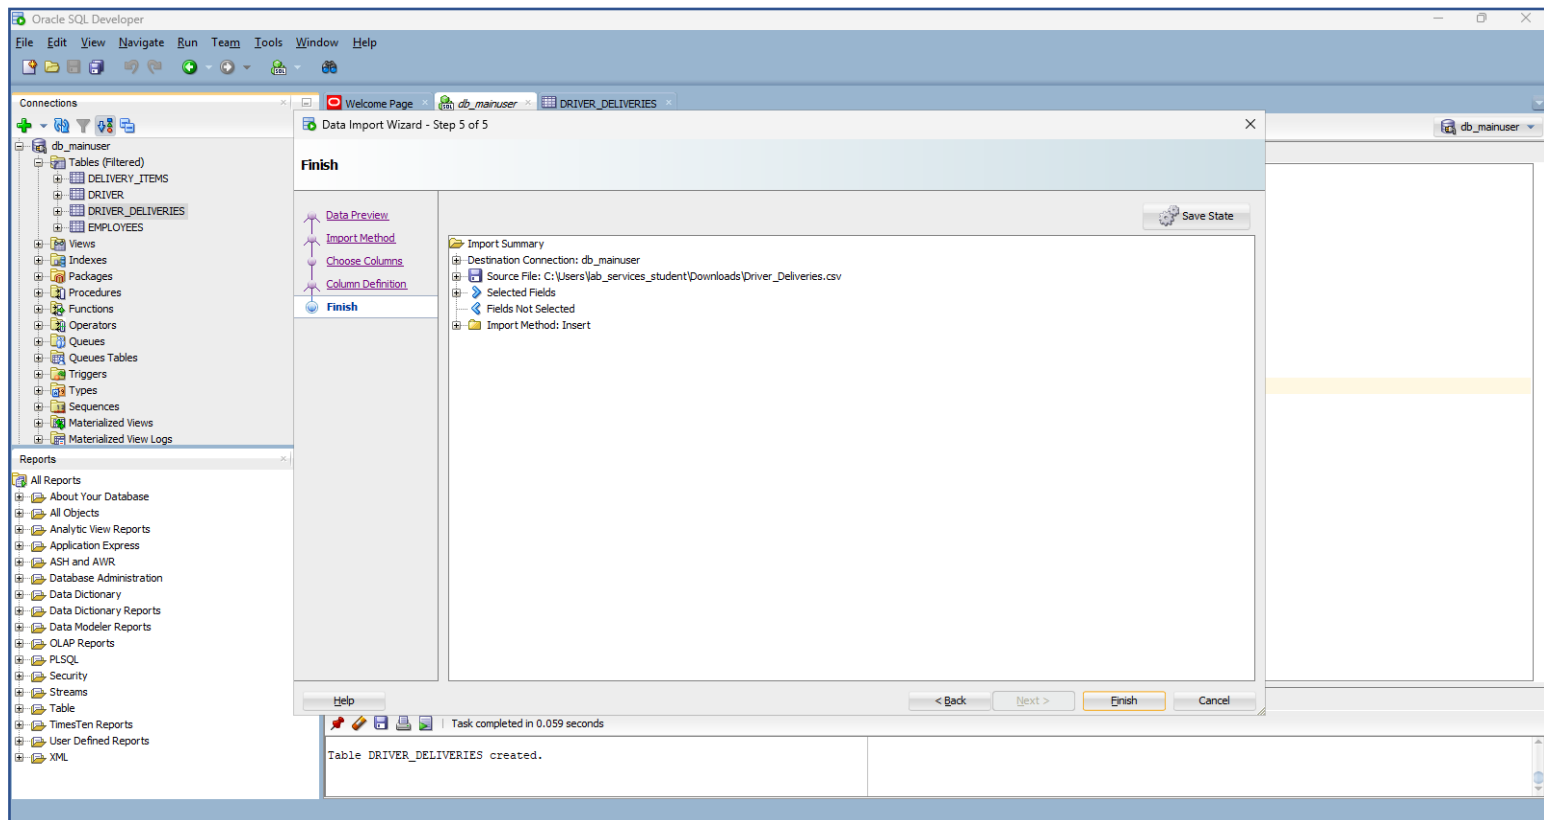
Question 1

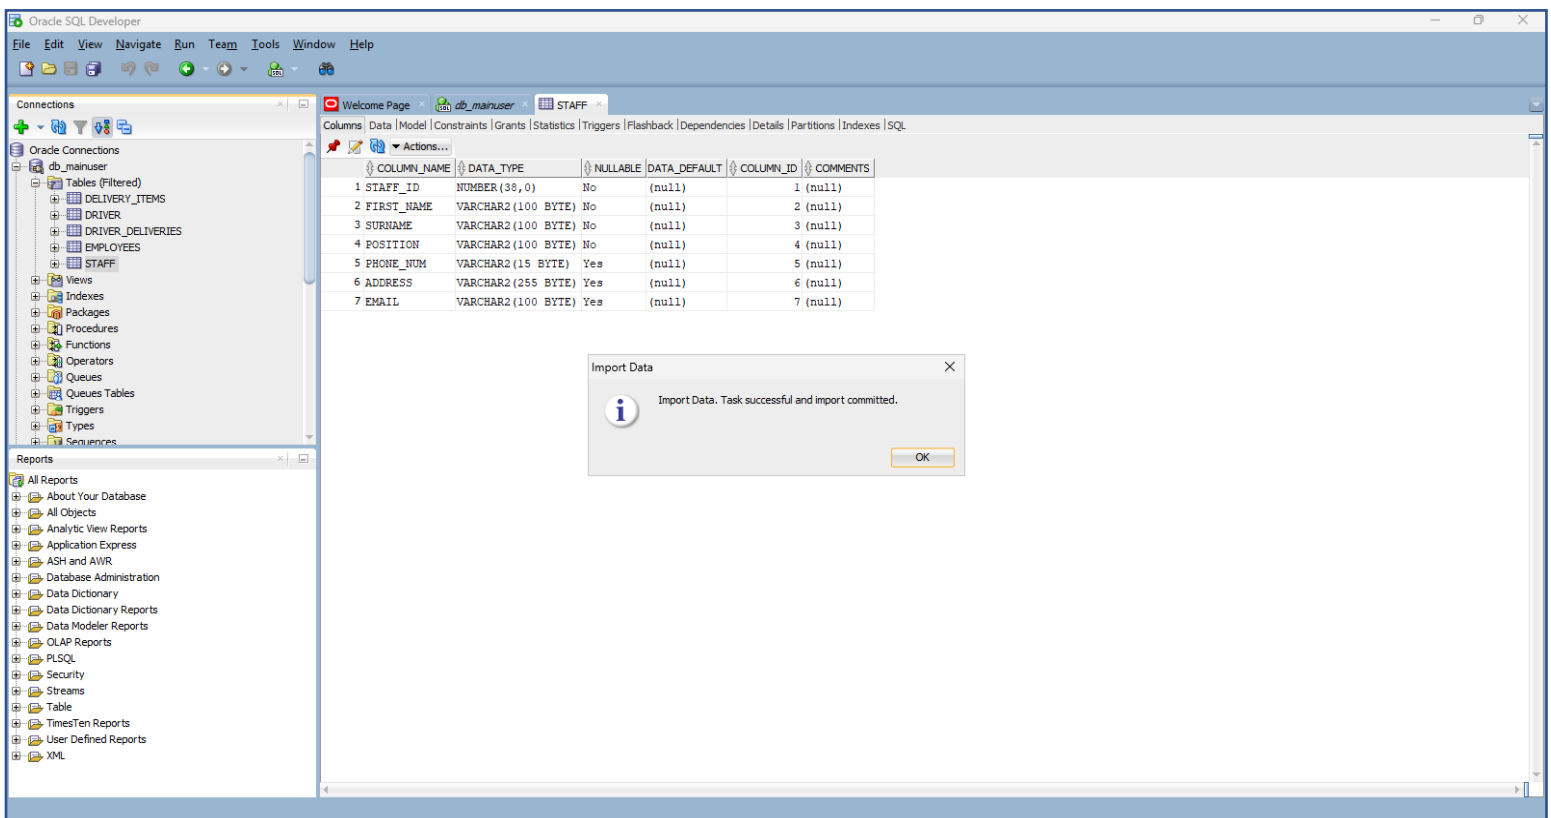
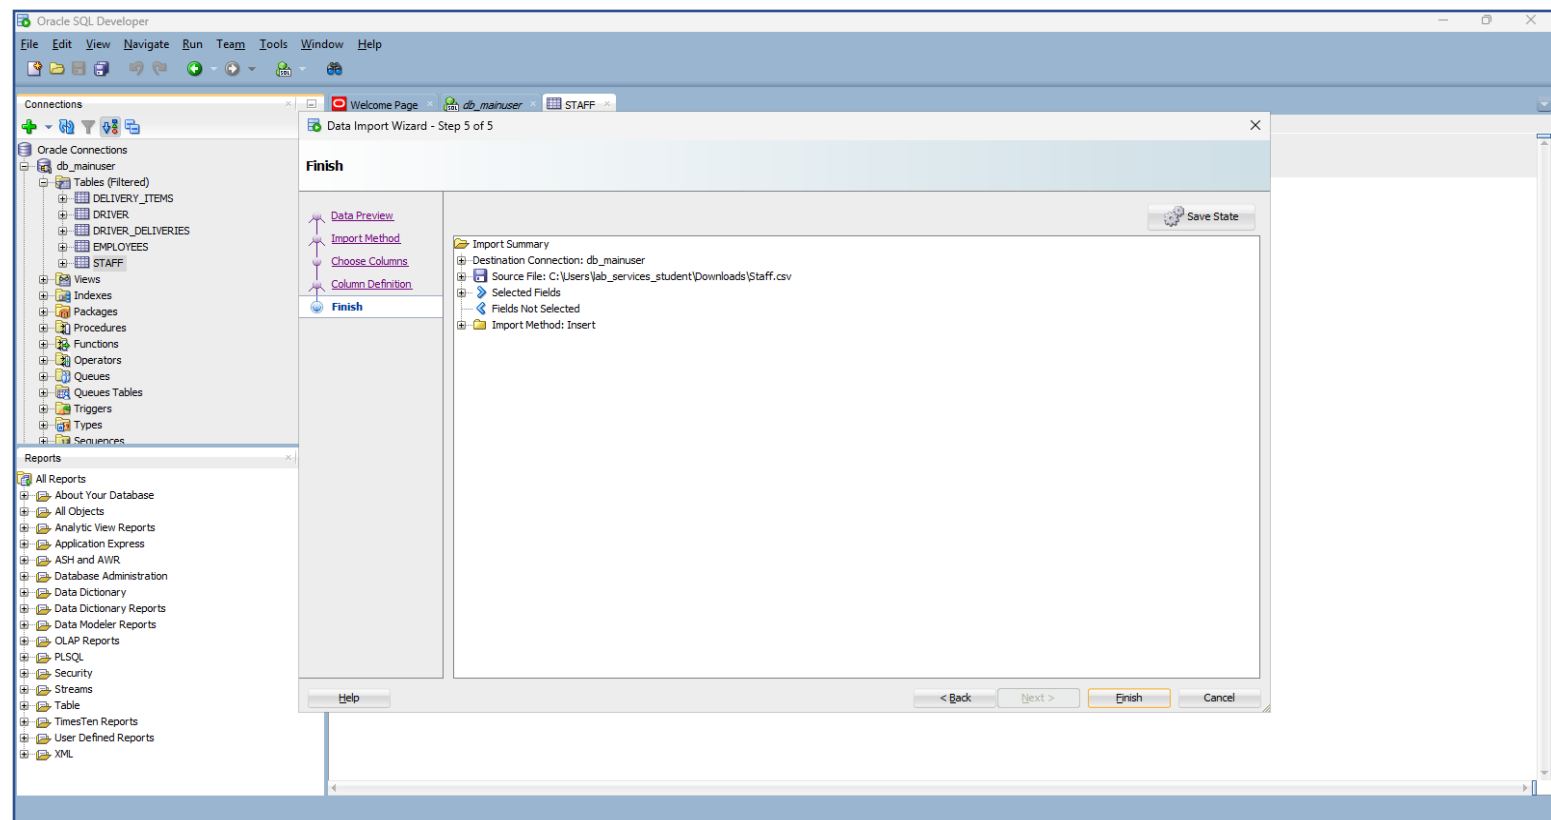


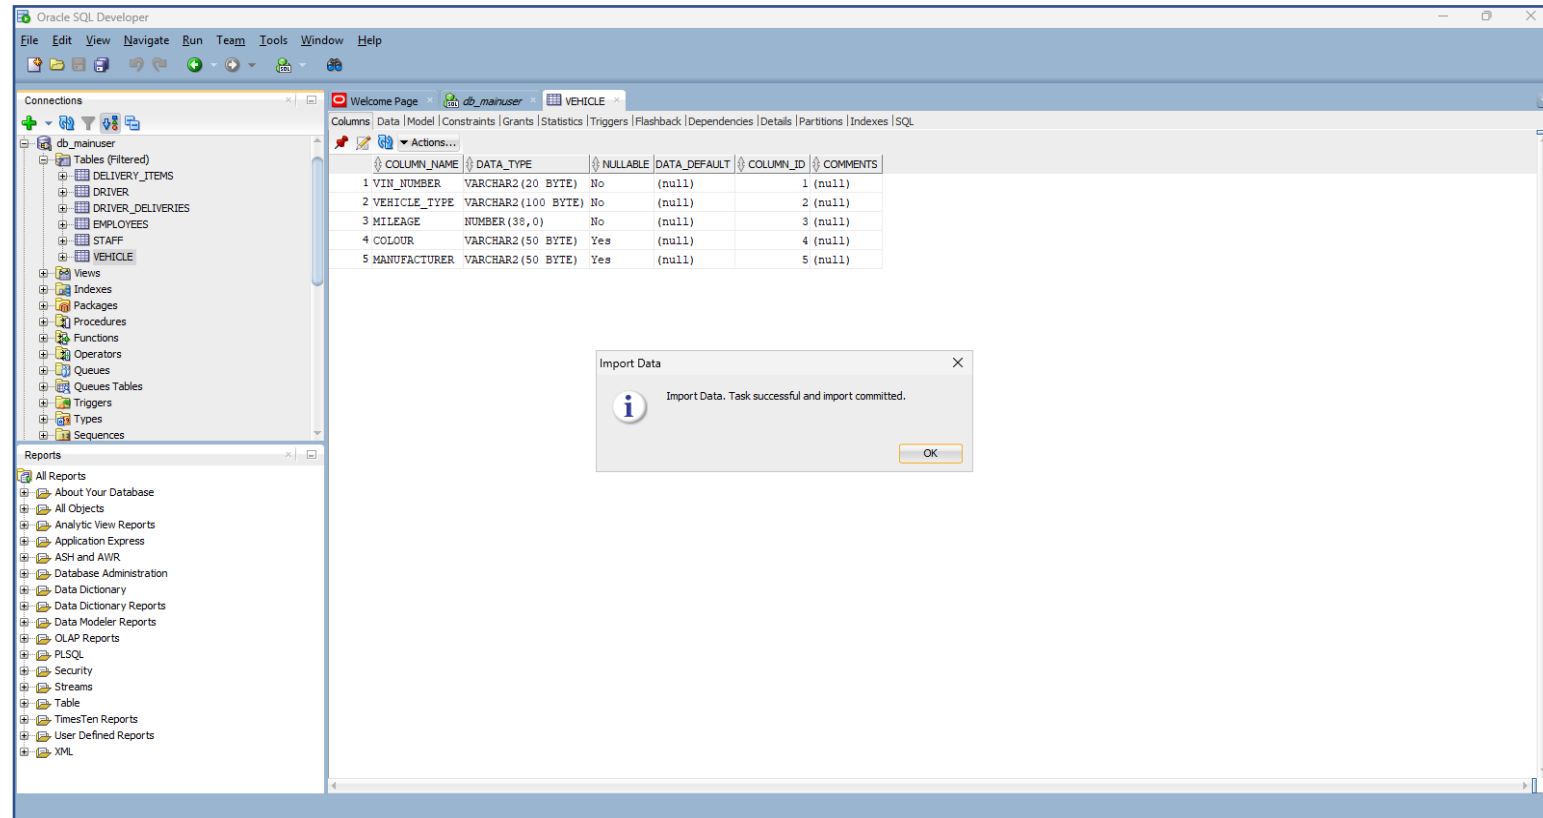

Question 2

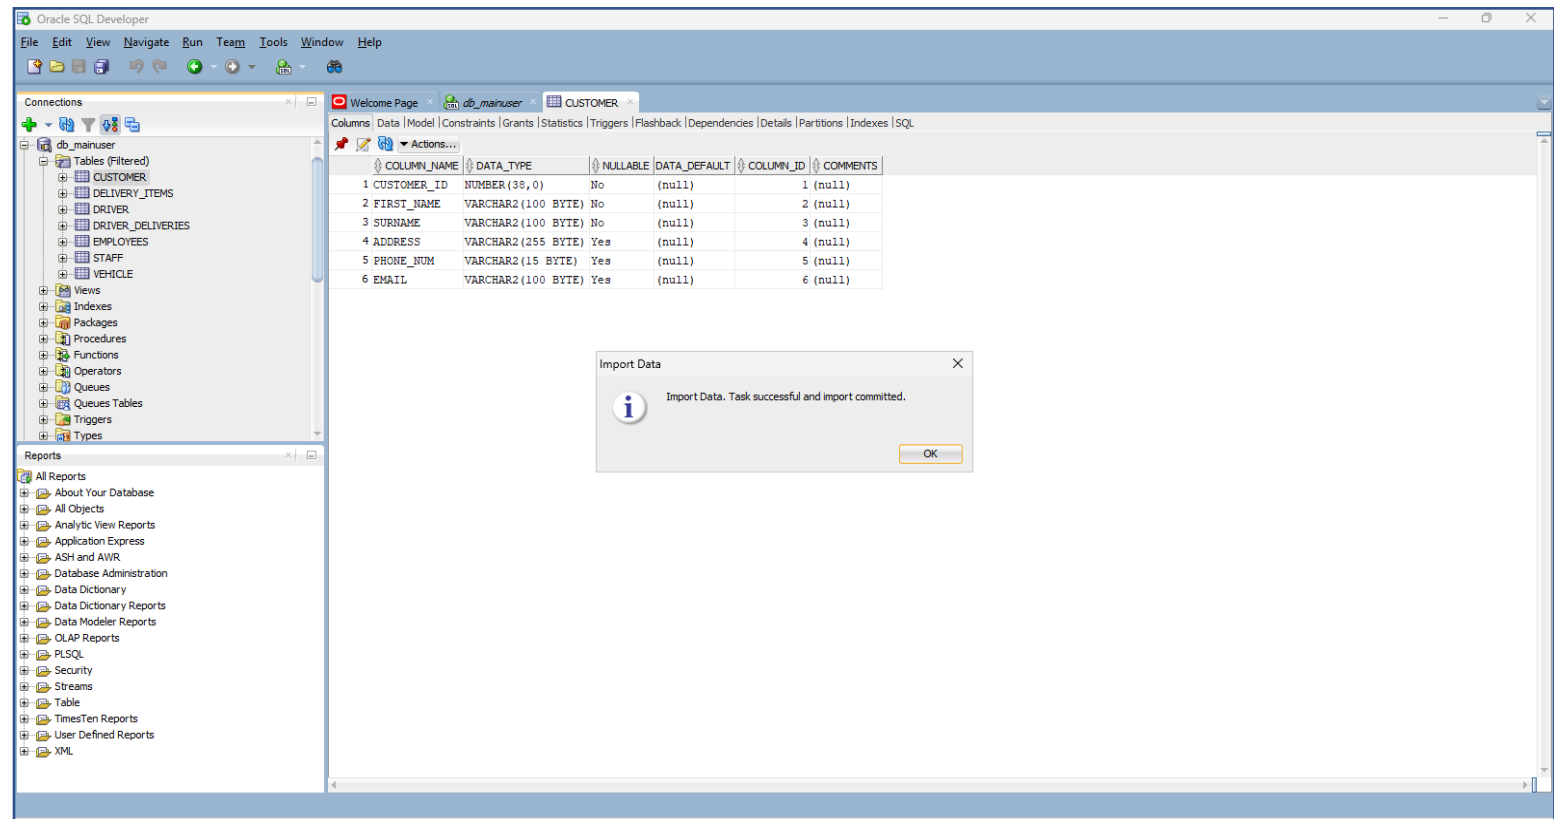
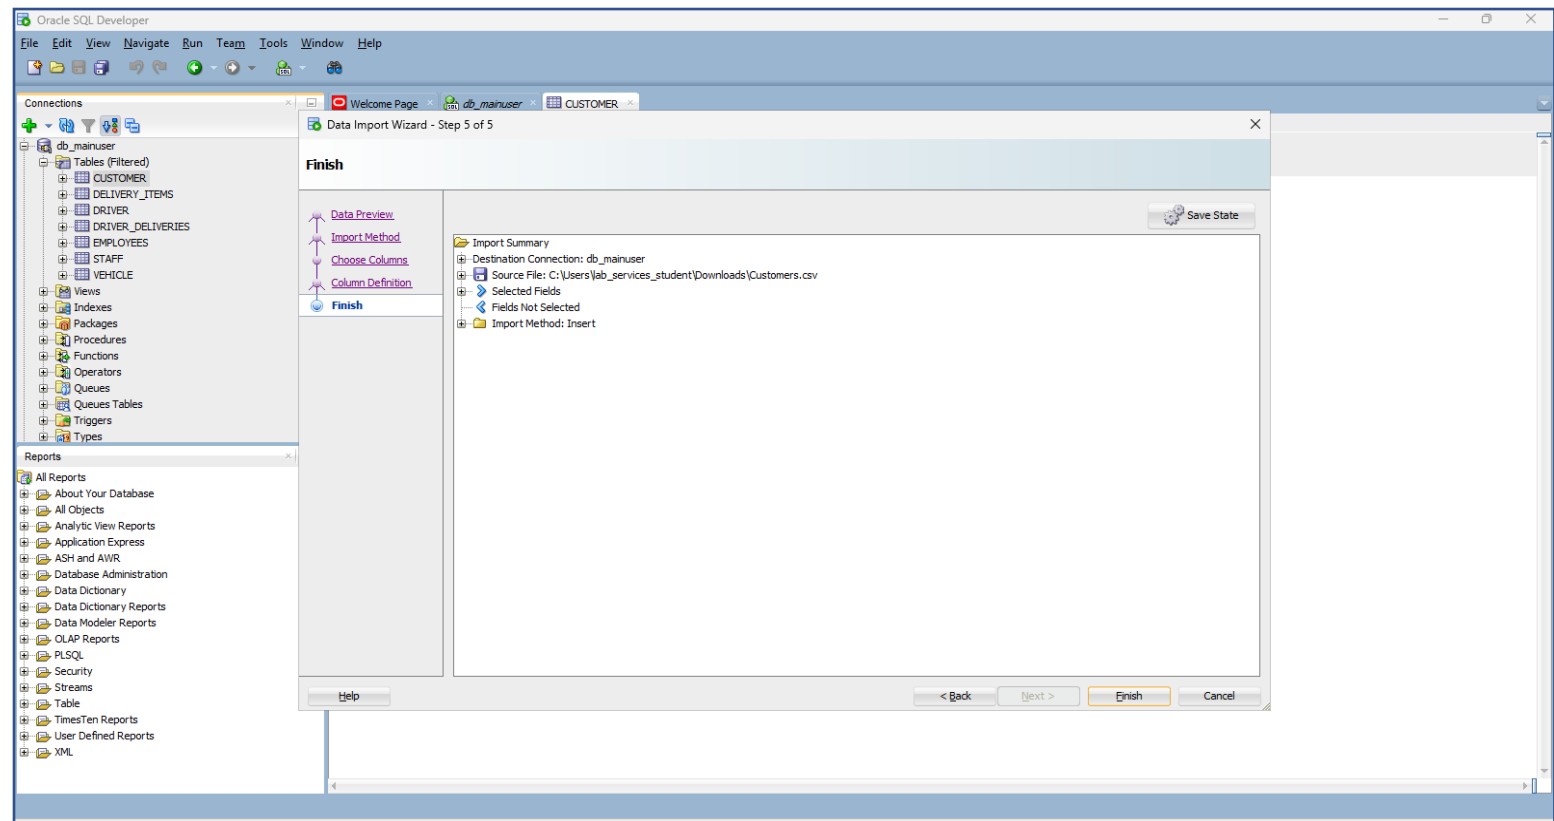



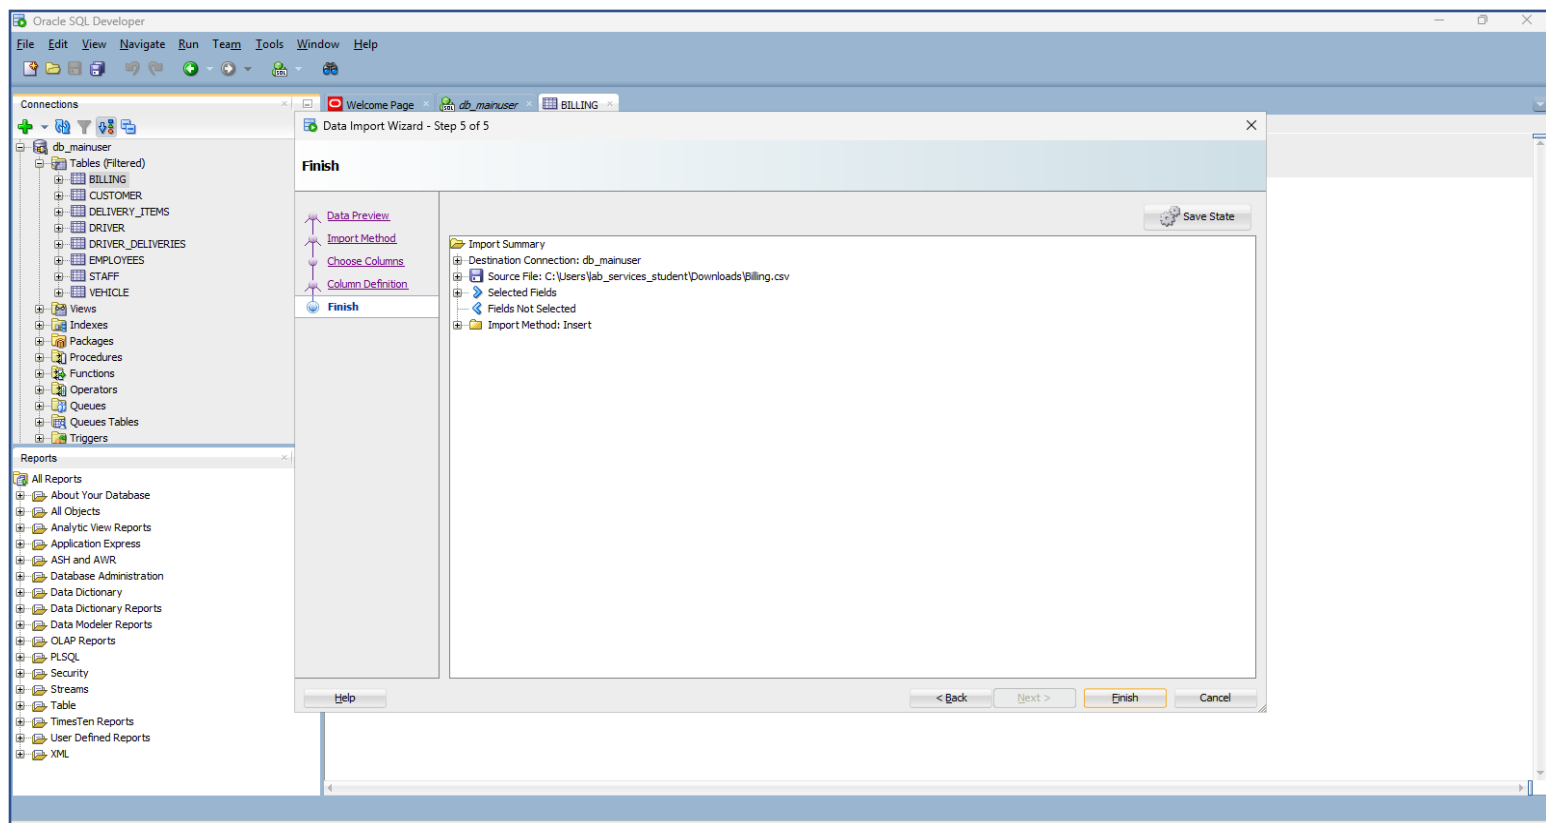












```
CREATE USER John IDENTIFIED BY Johnch2024;  
GRANT CREATE SESSION TO John;  
GRANT INSERT ANY TABLE TO John;
```

Script Output x

Task completed in 0.043 seconds

User JOHN created.

Grant succeeded.

Grant succeeded.

```
CREATE USER John IDENTIFIED BY Johnch2024;  
GRANT CREATE SESSION TO John;  
GRANT Select ANY TABLE TO John;
```

```
CREATE USER Hannah IDENTIFIED BY Hannahch2024;  
GRANT CREATE SESSION TO Hannah;  
GRANT INSERT ANY TABLE TO Hannah;
```

Script Output x

Task completed in 0.091 seconds

User HANNAH created.

Grant succeeded.

Grant succeeded.

Question 3.2

Improves the Security of Data:

- Only those who require it can access sensitive data.
- John is able to report on or analyse data, but he cannot alter it.
- Hannah is able to add new records without being able to see private data.

Reduces Unintentional Damage

- John is unable to inadvertently add or remove data while conducting queries.
- When Hannah inserts records, she cannot unintentionally reveal private information.

Question 4

The screenshot displays the Oracle SQL Developer interface. The left-hand pane shows the 'Connections' tree with 'db_mainuser' selected, and the 'Reports' tree below it. The main workspace is titled 'Worksheet' and contains the following PL/SQL code:

```
CREATE USER Hannah IDENTIFIED BY Hannahch2024;
GRANT CREATE SESSION TO Hannah;
GRANT INSERT ANY TABLE TO Hannah;

SET SERVEROUTPUT ON;

BEGIN
  FOR rec IN (
    SELECT d.First_Name, d.Surname, d.Driver_Code, v.VIN_Number, v.Mileage
    FROM Driver d
    JOIN Driver_Deliveries dd ON d.Driver_ID = dd.Driver_ID
    JOIN Vehicle v ON dd.VIN_Number = v.VIN_Number
    WHERE v.Mileage < 80000
  ) LOOP
    DBMS_OUTPUT.PUT_LINE('DRIVER: ' || rec.First_Name || ', ' || rec.Surname);
    DBMS_OUTPUT.PUT_LINE('CODE: ' || rec.Driver_Code);
    DBMS_OUTPUT.PUT_LINE('VIN NUMBER: ' || rec.VIN_Number);
    DBMS_OUTPUT.PUT_LINE('MILEAGE: ' || rec.Mileage);
    DBMS_OUTPUT.PUT_LINE(''); -- blank line between records
  END LOOP;
END;
```

The 'Script Output' pane at the bottom shows the results of the procedure execution:

```
DRIVER: Jomo, Mvuyisi
CODE: EC1
VIN NUMBER: 12A35868540
MILEAGE: 79058

PL/SQL procedure successfully completed.
```

The status bar at the bottom indicates 'Task completed in 0.075 seconds' and 'Line 91 Column 1 | Insert | Windows: C'.

Question 4.2

All of the data in a flat file database is kept in a single table or file, which frequently results in data redundancy, makes it harder to enforce data integrity, restricts query capabilities, and makes it difficult to scale. A relational database, on the other hand, uses primary and foreign keys to arrange data into several tables with precisely defined relationships, eliminating duplication and guaranteeing accuracy. Because it can effectively manage numerous entities, including drivers, vehicles, delivery items, staff, customers, and billing, a relational model is more appropriate for CHEETAH DELIVERIES. It can grow with the company, guarantees data integrity, and enables robust SQL queries for reporting. Management can track deliveries, create accurate reports, and keep consistent information without duplicating data by dividing entities into linked tables, which makes the relational model the best option over a flat file approach. (Staff, 2024)

Question 5

The screenshot displays the Oracle SQL Developer interface. The main window shows a SQL worksheet with the following query:

```
JOIN Delivery_Items d ON s.Staff_ID = d.STAFF_ID
GROUP BY s.Staff_ID, s.First_Name, s.Surname
ORDER BY delivery_count DESC
)
WHERE ROWNUM = 1; -- Pick the top staff member

-- Output exactly like the sample
DBMS_OUTPUT.PUT_LINE('STAFF ID: ' || v_staff_id);
DBMS_OUTPUT.PUT_LINE('FIRST NAME: ' || v_first_name);
DBMS_OUTPUT.PUT_LINE('SURNAME: ' || v_surname);
DBMS_OUTPUT.PUT_LINE('DELIVERIES PROCESSED: ' || v_deliveries);
END;
```

The bottom pane shows the script output:

```
Task completed in 0.258 seconds

STAFF ID: 51011
FIRST NAME: Sally
SURNAME: Du Toit
DELIVERIES PROCESSED: 1

PL/SQL procedure successfully completed.
```

The left sidebar shows the database structure, including tables, views, indexes, and other objects. The bottom status bar indicates the current line and column: "Line 120 Column 1 | Insert | Windows: C".

Question 5.2

The three primary parts of a PL/SQL block are Exception Handling, Execution, and Declaration.

Declaration Section: Variables, constants, cursors, and other components that the block will use are defined in this optional section. If we created a variable in the Q.5.1 query to hold the number of deliveries made by each employee, it would fall under the declaration portion. (Learning, N/A)

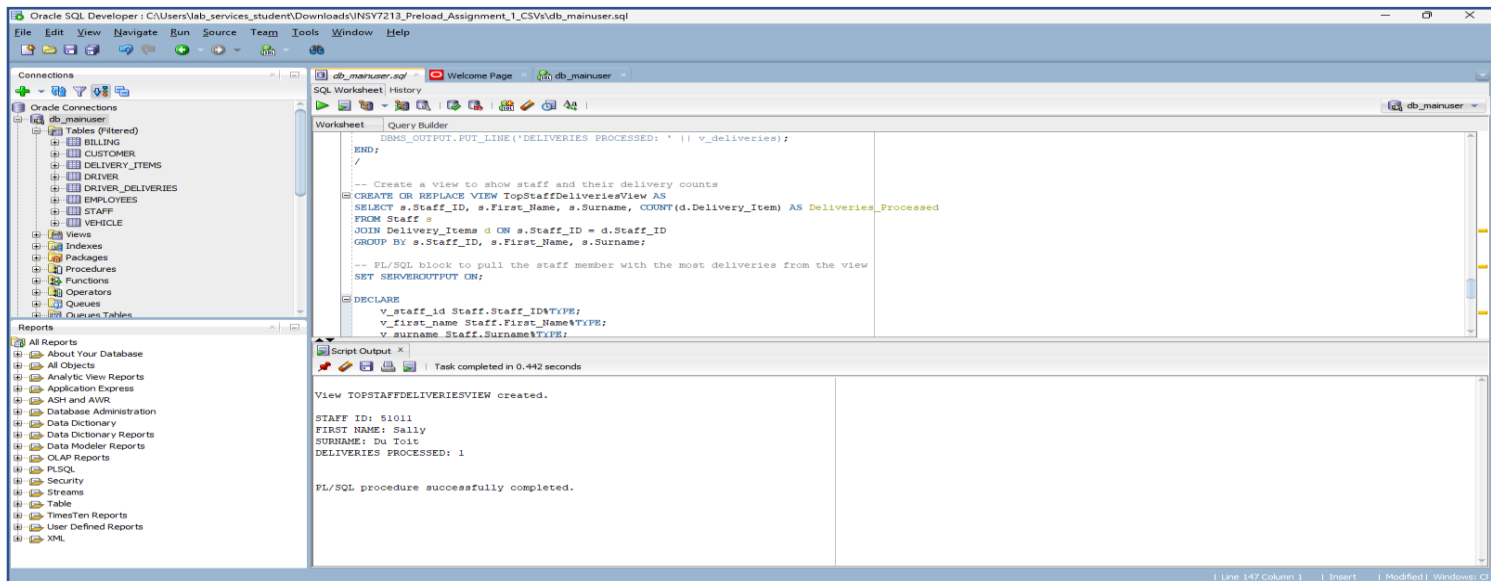
Execution part: The actual SQL statements and procedural logic that carry out the actions are contained in this required part. The SQL query that determines which employee has made the most deliveries is located in the execution section of Q.5.1. Here, the SELECT statement determines the top employees and counts deliveries. (Learning, N/A)

Exception Section: Any errors that arise during execution are handled in this optional section. For instance, the exception section can identify the error and show a helpful message rather than crashing if the query fails due to a missing table or division by zero in computations. (Learning, N/A)

Question 5.3.1

A view is a database virtual table that has a pre-written query. Instead of storing data, it dynamically shows the query's results whenever they are accessed. Management can fetch reports using a view instead than creating their own SQL queries. The advantage is that users can securely and regularly retrieve the same report without changing the underlying tables, which lowers the possibility of unintentional data changes. (Navlani, 2025)

Question 5.3.2



Question 6

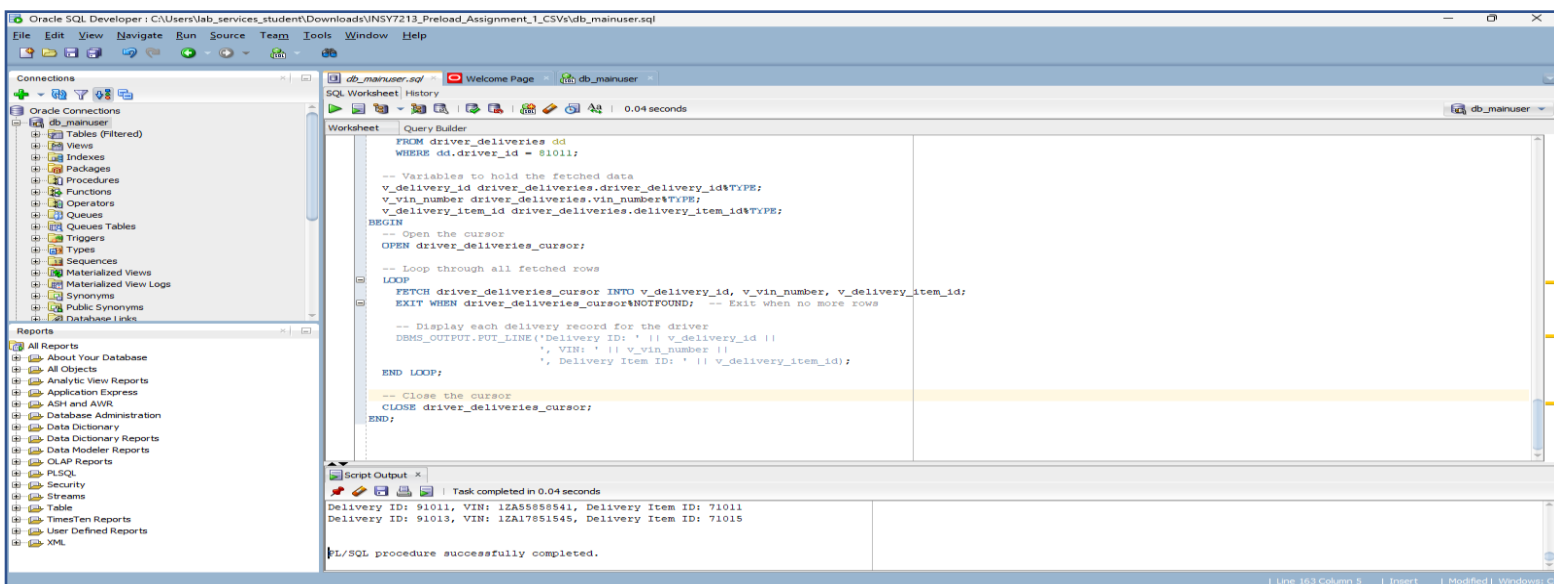
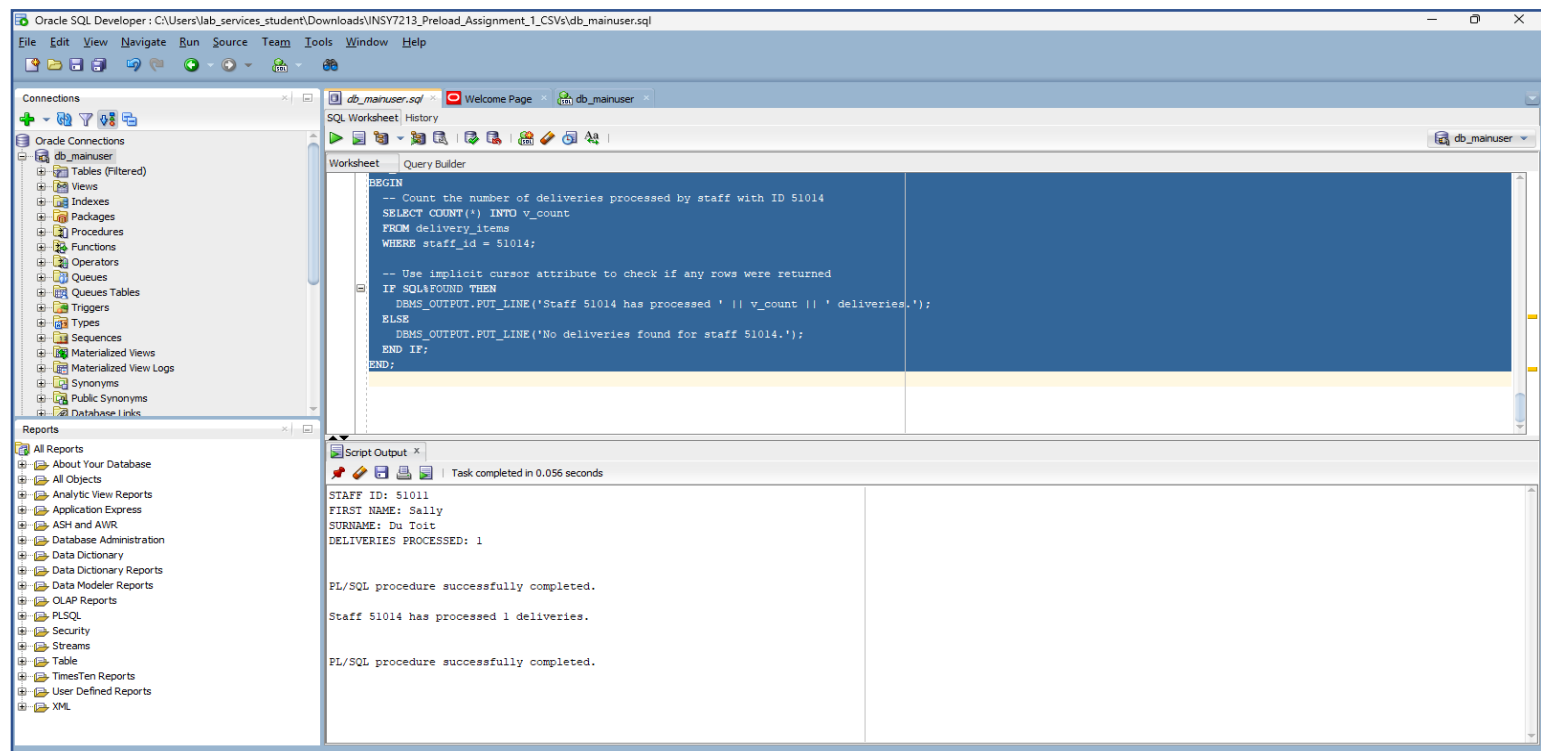
Question 6.1

Implicit Cursors

Every time a SQL query that modifies data is run, such as a SELECT INTO, INSERT, UPDATE, or DELETE statement, Oracle automatically creates an implicit cursor. Its objective is to give instant feedback regarding the SQL statement's execution without requiring manual cursor declaration or management. SQL%ROWCOUNT (number of rows impacted), SQL%NOTFOUND (TRUE if no rows were affected), SQL%FOUND (returns TRUE if the statement affected at least one row), and SQL%ISOPEN (always FALSE because implicit cursors are automatically opened and closed) are common attributes of implicit cursors. The simplicity of implicit cursors is what drives their use; they are perfect for DML operations or single-row queries where the developer wants to rapidly determine whether a statement was successful or how many rows were impacted without having to deal with the additional burden of opening, fetching, and closing a cursor by hand. For instance, in CHEETAH DELIVERIES, an implicit cursor can be used to instantly ascertain whether any deliveries exist and tally the number of deliveries processed by a staff member. (point, N/A)

Explicit Cursors

In order to get many rows from a query, an explicit cursor needs to be manually declared and managed by the developer. When working with several rows, it gives you more control than implicit cursors by enabling row-by-row processing of query results. `cursor_name%FOUND` (TRUE if the last fetch returned a row), `cursor_name%NOTFOUND` (TRUE if the last fetch did not return a row), `cursor_name%ROWCOUNT` (number of rows fetched thus far), and `cursor_name%ISOPEN` (TRUE if the cursor is currently open) are common properties of explicit cursors. Explicit cursors are used because they enable in-depth processing of every single row from a query, which is very helpful for creating reports or carrying out intricate operations on several records. To give management a comprehensive picture of driver activity, for example, in CHEETAH DELIVERIES, an explicit cursor can be used to loop over all deliveries for a particular driver and display each delivery's ID, vehicle VIN, and delivery item. (point, N/A)



Question 6.2

Sequences, which are commonly used for primary keys or other unique identifiers, are database objects that automatically create distinct number values. Its goal is to guarantee that every new row added to a database is assigned a distinct number without the need for manual computation or monitoring of the most recent value. The following are common sequence attributes: INCREMENT BY (the amount to increment for each new value), NOCYCLE or CYCLE (if the sequence restarts after reaching its maximum value), CACHE or NOCACHE (whether values are preallocated for quicker performance), and START WITH (the first number created). Sequences are used to ensure uniqueness, avoid conflicts in multi-user setups, and make record entry easier. To ensure that every delivery item has a unique identity without having to manually check existing IDs, for instance, CHEETAH DELIVERIES can utilise a sequence to produce DELIVERY_ITEM_ID automatically when a new delivery is added. This enhances efficiency and preserves data integrity. (oracle, N/A)

The screenshot displays the Oracle SQL Developer interface. The main window shows a SQL script with the following content:

```
LOOP
  FETCH driver_deliveries_cursor INTO v_delivery_id, v_vin_number, v_delivery_item_id;
  EXIT WHEN driver_deliveries_cursor%NOTFOUND; -- Exit when no more rows

  -- Display each delivery record for the driver
  DBMS_OUTPUT.PUT_LINE('Delivery ID: ' || v_delivery_id ||
    ', VIN: ' || v_vin_number ||
    ', Delivery Item ID: ' || v_delivery_item_id);

END LOOP;

-- Close the cursor
CLOSE driver_deliveries_cursor;
END;
```

Below the script, there is a comment and a SQL statement to create a sequence and insert a row:

```
-- Create a sequence for BILL_ID
CREATE SEQUENCE bill_seq
  START WITH 1000
  INCREMENT BY 1
  NOCACHE;

-- Use the sequence when inserting a new bill
INSERT INTO Billing (Bill_ID, Customer_ID, Staff_ID, Bill_Date)
VALUES (bill_seq.NEXTVAL, 11011, 51014, SYSDATE);

-- Check result
SELECT * FROM Billing WHERE Bill_ID >= 1000;
```

The bottom of the window shows the 'Query Result' tab with the following data:

BILL_ID	CUSTOMER_ID	STAFF_ID	BILL_DATE
1	1000	11011	51014 26-SEP-25

References

Khusid, A., 2011. *Miro*. [Online]

Available at: <https://miro.com>

[Accessed 20 September 2025].

Learning, G., N/A. *Great Learning*. [Online]

Available at: <https://www.mygreatlearning.com/plsql/tutorials/parts-of-pl-sql-subprogram>

[Accessed 20 September 2025].

Navlani, A., 2025. *Data Camp*. [Online]

Available at: <https://www.datacamp.com/tutorial/views-in-sql>

[Accessed 20 September 2025].

oracle, N/A. *Oracle*. [Online]

Available at: <https://docs.oracle.com/en/database/oracle/oracle-database/19/sqlrf/CREATE-SEQUENCE.html>

[Accessed 20 September 2025].

point, T., N/A. *Tutorials Point*. [Online]

Available at: https://www.tutorialspoint.com/plsql/plsql_cursors.htm

[Accessed 20 September 2025].

Staff, E., 2024. *ECL*. [Online]

Available at: <https://www.ecisolutions.com/blog/building-supply/the-difference-between-flat-file-and-relational-data/>

[Accessed 20 September 2025].